

COMP2002

ARTIFICIAL INTELLIGENCE

20 CREDIT MODULE

ASSESSMENT: 100% Coursework

W1: 30% Set Exercises

W2: 70% Report

MODULE LEADER: Dr David Walker

MODULE AIMS

- To familiarise with the underlying principles of artificial intelligence.
- To expose students to topics such as search and optimization, knowledge representation and reasoning, and machine learning.
- To instil students with an appreciation of the importance of ethical considerations behind artificial intelligence.

ASSESSED LEARNING OUTCOMES (ALO):

1. Describe and analyse a range of artificial intelligence methods and their applications.
2. Compare artificial intelligence paradigms and evaluate the appropriateness of a particular paradigm for specific application domains.
3. Choose and apply appropriate artificial intelligence methods to a chosen sample domain.

Overview

This document contains all the necessary information pertaining to the assessment of *COMP2002 Artificial Intelligence*. The module is assessed via **100% coursework**, across two elements: 30% *Set Exercises* and 70% *Report*.

The sections that follow will detail the assessment tasks that are to be undertaken. The submission and expected feedback dates are presented in Table 1. All assessments are to be submitted electronically via the respective DLE module pages before the stated deadlines.

	Submission Deadline	Feedback
Set Exercises (30%)		
SE1 (30%)	2nd March 2021, 12pm	30th March 2021
SE2 (40%)	20th April 2021, 12pm	19th May 2021
SE3 (30%)	25th May 2021, 12pm	16th June 2021
Report (70%)	19th May 2021, 12pm	16th June 2021

Table 1: Assessment Deadlines

All assessments will be introduced in class to provide further clarity over what is expected and how you can access support and formative feedback prior to submission. Whilst the assessment information is provided at the start of the module, it is not necessarily expected you will start this immediately – as you will often not have sufficient understanding of the topic. The module leader will provide guidance in this respect.

Set Exercises

There are three set exercises for you to complete during the course of the module. Each of these exercises should take you approximately 8 hours to complete and you have four weeks to complete them.

Assessment 1: Set Exercises

Set Exercise 1: worth 30% of element mark

During Lecture 1 you were asked to suggest applications of artificial intelligence. Your task is to **select one of these applications** and **write a blog post describing it**.

Your blog post must:

- Introduce the topic – you should explain the purpose behind the AI. What will it mean we can now do that we couldn't before?
- Discuss whether the AI has been successful – what has the impact of its introduction been.
- Discuss the future of the AI – will it continue to be used? Are there any improvements required to make it better?

The word limit for your blog post is **500 words**.

For information on how to write a good blog post, have a look at the resources posted on the DLE.

You must **submit your blog post in PDF form**.

Assessment Criteria

A rubric will be used to **assess** and **provide feedback** on your submission. This is provided below:

AI application is introduced	0 points Requirements not met	9 points Vague introduction with little detail	11 points Some introduction but some omissions	13 points Application idea introduce but more detail on its purpose needed	15 points Generally good introduction – AI and its purpose are discussed. Some minor omissions.	20 points The AI and its purpose are fully introduced. No omissions.
The AI's success is discussed	0 points Requirements not met	11 points Vague description of AI's impact with little detail. No discussion.	17 points Some description of AI's impact in some detail. No discussion.	20 points Good description of the AI's impact. Some analytical discussion.	23 points The AI's impact is well described. Good analytical discussion but lacking detail in a few places.	30 points Excellent discussion of the AI's impact and strong analytical discussion.
The AI's future is discussed	0 points Requirements not met	11 points Vague description of the AI's future with little detail. No discussion.	17 points Some description of the AI's future in some detail. No discussion.	20 points Good description of the AI's future. Some analytical discussion.	23 points The AI's future is well described. Good analytical discussion but lacking detail in a few places.	30 points Excellent discussion of the AI's future and strong analytical discussion.
Style and language	0 points Requirements not met	9 points Poor spelling and grammar. Inappropriate structure.	11 points Poor spelling and grammar. Some improvements to structure required.	13 points Some minor spelling or grammatical errors. Some improvements to structure required.	15 points Some minor spelling or grammatical errors. A suitable structure is followed.	20 points Excellent. Very few spelling or grammatical errors. A suitable structure is followed.

Assessment 1: Set Exercises

Set Exercise 2: worth 40% of element mark

Your task for this set exercise is to write a **1,000 word** (+/- 10%) essay on the title “**bias will prevent the widespread adoption of AI**”. You must produce an argument that either agrees or disagrees with the statement, and provide **evidence** to back up your claims.

Your essay should:

- Introduce the topic – explain what bias is and how it might affect an AI.
- State why you either agree or disagree with the statement.
- Provide evidence for your agreement or disagreement.
- Conclude.

It is important that you discuss the material you are including – think analytically about it, don't just provide a series of facts. The idea is for you to take a position, but I don't mind what that position is.

This is an academic essay so its style should reflect that – you must include references. You should consider a range of sources (include peer-reviewed journal papers as well as other sources such as news sites). When including a source, consider its reliability. Can it be trusted? References must be presented using the **Harvard style** – more information on this is on the module DLE page – and should be provided as a list at the end of the essay and are not counted within the word count.

You must **submit your essay in PDF form**.

Please turn over for assessment criteria.

Assessment Criteria

A rubric will be used to **assess** and **provide feedback** on your submission. This is provided below:

Grammar and spelling	0 points Requirements not met	9 points Very poor spelling and grammar.	11 points Reasonable quality, though a number of spelling and grammar errors.	13 points Good presentation, with a number of small spelling or grammar errors.	15 points Clear presentation with a small number of spelling or grammar errors.	18 points Excellent. Few spelling or grammar errors could be found.	20 points Excellent. No spelling or grammar errors could be found.
References and citations	0 points Requirements not met.	9 points Very few sources of questionable reliability included. References are not presented correctly.	11 points Some reliable sources included though more are needed. References are not presented correctly.	13 points A reasonable range of sources are included – some are reliable. There are errors in their presentation.	15 points A reasonable range of mostly reliable sources included. There are a few issues with their presentation.	18 points A good range of mostly reliable sources included. References are mostly presented correctly.	20 points A good range of sources included – all reliable. All references are correctly presented.
Content	0 points Requirements not met	18 points Poorly structured concepts with no development of ideas and little evidence.	22 points Evidence of organisation of concepts but only adequately structured.	26 points Clear organisation of the concepts but not always logical or flowing.	30 points Concepts mostly well organised, generally well supported with evidence.	36 points Concepts well organised and mostly support with evidence.	40 points Concepts well organised and well supported with evidence.
Argument	0 points Requirements not met	9 points A position has been taken but is supported by a poor argument.	11 points A position is taken and there is some attempt at providing a supporting argument.	13 points The position of the author is clear and there is a weak supporting argument.	15 points The position of the author is clear and there is a supporting argument that needs reinforcing in places.	18 points The position of the author is clear and is supported by an excellent argument.	20 points The position of the author is clear and is supported by an excellent argument and balanced discussion.

Assessment 1: Set Exercises

Set Exercise 3: worth 30% of element mark

Your task here is to produce a technical poster about the design of an AI-based assistant to be used by the University at open days. You must outline the AI components that will go into the assistant, explain how a user will interact with it.

Your poster must:

- Describe the AI components – how will you use machine learning, evolutionary computation, knowledge representation and NLP? (You might not use all of these, it's up to you to decide). How and why will each be used?
- Indicate how the components are assembled into the final system.
- Outline the UI – how does the user input requests, and how is the information presented?

The poster must be submitted as a PDF file to the DLE and you may use a single page only. A template has been provided for you on the DLE.

For information on how to write a good scientific poster, have a look at the resources posted on the DLE.

Assessment Criteria

A rubric will be used to **assess** and **provide feedback** on your submission. This is provided below:

Selection and justification of AI components	0 points Requirements not met	20 points Not clear which AI will be used or how, no rationale behind their selection.	27 points Some of the AI's use is described but there is little rationale.	33 points The AI to be used is mostly described. Some aspects are justified but more detail is needed.	38 points The AI to be used is described. Most areas are justified but justification could be strengthened in places.	45 points The AI to be used is described and is mostly justified.	50 points All of the AI to be used is described and the rationale behind the choices is clear.
UI design	0 points Requirements not met	14 points Little consideration to the UI is given, and there is no rationale behind the design.	17 points There is some consideration to the UI and some reasonable choices have been made. No justification.	13 points Reasonable UI choices have been made and justification has been attempted.	20 points The UI is mostly well designed. The design is generally justified but more depth is needed.	27 points The UI is well designed. The rationale behind the design is mostly clear but could be enhanced in one or two areas.	30 points The UI is well designed and the rationale behind the design is clear.
Style and language	0 points Requirements not met	9 points Poor spelling and grammar. Inappropriate structure. Poster content is too brief.	11 points Some spelling and grammar errors. Inappropriate structure. Poster content is too brief.	13 points Some spelling and grammar errors. Minor improvements to structure required. Poster content is too brief.	15 points Some minor spelling or grammatical errors. Minor improvements to structure required.	18 points Some minor spelling or grammatical errors. A suitable structure is followed.	20 points Excellent. Very few spelling or grammatical errors. A suitable structure is followed.

Assessment 2: Machine Learning and Optimisation

This assignment contributes **70%** of the overall module mark for COMP2002 and is an **individual assignment**. You must submit the deliverables to the DLE by the specified submission dates.

The coursework has two parts – one is a **machine learning exercise** and the second is about **evolutionary computation**. You must **complete and submit both parts**. Each part is worth 50% of the coursework mark. A Jupyter notebook has been placed on the DLE for you to use. You should download it and use it to implement the code you need to complete the tasks below.

PART 1 – MACHINE LEARNING

You have been provided with datasets relating to the potential for forest fires in two regions in Algeria. Your task is to train regression models that predict a fire likelihood index based on 12 numerical inputs.

You must complete the following tasks:

Task 1.1 – Data preparation (10% of total mark)

The first phase of the work requires you to load the data you have been provided with into your Python program. Before the data can be used to train and test your models you must first prepare it – this means that the inputs must be normalized. There is no missing data in the dataset.

Task 1.2 – Regression (20% of total mark)

Having prepared the data you must now build a regression tool that can predict new points. Use the following regression implementations within the *scikit-learn* package to construct predictors for the dataset:

- Random Forest (*sklearn.ensemble.RandomForestRegressor*)
- Neural Network (*sklearn.neural_network.MLPRegressor*)
- Support Vector Machine (*sklearn.svm.SVR*)

You must demonstrate that each regressor is capable of providing a prediction for a given input.

Task 1.3 – Assessment of regression (20% of total mark)

The regression models you have used in the previous task must be assessed. To do this you are required to assess the *mean absolute error* rate for each model. You may use the MAE implementation available in *scikit-learn* to do this. It is not sufficient to report a single MAE rate. You must use cross validation to report training and testing results and report these values using a boxplot.

PART 2 – EVOLUTIONARY COMPUTATION

The second part of this assignment requires you to implement an evolutionary algorithm to solve the travelling salesperson problem (TSP) for a pair of routes. You have been provided with a pair of files, each of which describes the distance between cities. Your task is to optimise a route so that each city is visited once by following the shortest possible route.

Task 2.1 – Generation of random solutions (10% of total mark)

Your first task is to implement the fitness function. Your code should read one of the files and given a random route it should return the distance the salesperson would travel following it. You should call your function and print out the distance travelled for the random route.

Task 2.2 – Algorithm implementation (25% of total mark)

You should implement a hillclimber (as described in the lectures) to optimise both TSP instances problem. Your algorithm must have the following features:

- It should be possible to use one of two mutation operators – one is a **swap** operator (taking two cities in the route and swapping them) and the other is a **ruin-and-recreate** operator (which generates a completely new solution at random).
- At each iteration your single **parent** solution should be used to generate a single **child** operator by using **one** of the mutation operators (either the swap or the ruin-and-recreate for the entire run of the algorithm).
- At the end of an iteration the algorithm should retain the parent or child that has the best fitness.
- At the end of an iteration the best (parent or child) solution's fitness should be added to a list of the best fitnesses, which is returned along with the best solution at the end of the optimisation.

Task 2.3 – Visualisation of results (15% of total mark)

For both TSP instances you should run the algorithm twice – once for each mutation operator. Repeat this 30 times, so that you get 30 fitness lists for the swap operator and 30 fitness lists for the ruin-and-recreate operator. Plot the **average**, **maximum** and **minimum** fitness at each iteration for each operator. You should plot them on the same graph so that they can be compared. You should be able to see which optimiser is best – state which in the notebook, and say why.

COURSEWORK DELIVERABLES

A Jupyter notebook has been provided on the DLE for you to use for this coursework. You should implement your code in it, and submit it to the DLE ahead of the deadline specified in the submission dates earlier in this document. **Please indicate which task each section of the notebook refers to using a Markdown cell.**

Please check your submitted files are correct by downloading them again and checking that they work. **You will receive a confirmation receipt by email when your work has been properly submitted** – if you do not receive this email then your work has not been submitted.

ASSESSMENT CRITERIA

Your work will be assessed according to the rubric found in Tables 2 and 3. Your mark for this piece of coursework will be based on an aggregation of the marks for each category. Marks will be awarded based on **both the demonstration and the report**.

Category	Fail (< 40%)	>= 40%	>= 50%	>= 60%	>= 70%	
Data preparation (10%)	Data is badly loaded into the program and there is no real preparation of the data.	The data is loaded correctly, though inefficiently, but there is no preparation.	The data is efficiently loaded into the program and there is an attempt at normalisation.	The data is loaded and normalised, with some inefficiencies.	The data is loaded and normalised, and the code to do so is efficiently written.	/10
Regression (20%)	Very little or no attempt at training or testing regression models. Poor organisation of code.	There has been an attempt at training one or more of the models, but there is no testing. The organisation of the code is poor.	All three models are trained, but there are problems with the testing. The organisation of the code is poor.	All three models are trained and tested correctly. The code to do so is repetitive and not well organised.	All three models are trained and tested correctly. The code used to do so is well organised and efficient.	/20
Assessment of results (20%)	Very little or no attempt at assessing the regression results.	There has been an attempt at using mean absolute error, but little or no use of cross validation.	Mean absolute error is used in combination with cross validation. The presentation of results should be improved.	Mean absolute error is used in combination with cross validation, and the results are presented with a boxplot. Some inefficiencies in the code organisation.	Mean absolute error is used in combination with cross validation, and the results are presented with a boxplot. The code is efficient and well organised.	/20

Table 2: Feedback Template for Assessment 2 (Machine Learning part)

Category	Fail (< 40%)	>= 40%	>= 50%	>= 60%	>= 70%	
Generation of random solutions (10%)	The fitness function is incorrect. Poor attempt at loading the data. Generation of random solutions does not work.	The data is loaded. There has been an attempt at the fitness function, but it is incorrect. Generation of random solutions does not work.	The data is loaded and the fitness function is close to correct. Generation of random solutions does not work.	The data is loaded and the fitness function is close to correct. Solution generation is close to working. Code is inefficiently designed.	The data is loaded and the fitness function is correct. Random solutions are generated correctly. The code to do so is efficiently structured.	/10
Algorithm implementation (25%)	Mutation operators are mostly incorrect or missing. The algorithm is incorrect – selection is missing or not working and the fitness archive is missing.	One mutation operator has been attempted. Some of the algorithm is present, but it is incomplete or incorrect. The fitness archive has been attempted.	Two mutation operators have been attempted but are incomplete or incorrect. The algorithm is partially correct and the fitness archive works correctly.	The mutation operators work and the algorithm is mostly implemented correctly. The fitness archive stores fitnesses correctly.	The mutation operators and algorithm have been implemented correctly. The fitness archive stores fitnesses correctly. The code is efficiently structured and well organised.	/25
Visualisation of results (15%)	No attempt at producing a plot, or the plot is incorrect. Experiments are not run for the correct number of repeats.	A graph shows the partial results required. The experiment was run incorrectly.	A graph shows the correct results required but the experiment was run incorrectly.	A graph shows the correct results and the correct experimental setup has been followed. No analysis of the results.	A graph shows the correct results and the correct experimental setup has been followed. The best mutation operator is identified and justified.	/15

Table 3: Feedback Template for Assessment 2 (Evolutionary Computation part)

General Guidance

Extenuating Circumstances

There may be a time during this module where you experience a serious situation which has a significant impact on your ability to complete the assessments. The definition of these can be found in the University Policy on Extenuating Circumstances here:

https://www.plymouth.ac.uk/uploads/production/document/path/15/15317/Extenuating_Circumstances_Policy_and_Procedures.pdf

Plagiarism

All of your work must be of your own words. You must use references for your sources, however you acquire them. Where you wish to use quotations, these must be a very minor part of your overall work.

To copy another person's work is viewed as plagiarism and is not allowed. Any issues of plagiarism and any form of academic dishonesty are treated very seriously. All your work must be your own and other sources must be identified as being theirs, not yours. The copying of another person's work could result in a penalty being invoked.

Further information on plagiarism policy can be found here:

Plagiarism: <https://www.plymouth.ac.uk/student-life/your-studies/essential-information/regulations/plagiarism>

Examination Offences: <https://www.plymouth.ac.uk/student-life/your-studies/essential-information/exams/exam-rules-and-regulations/examination-offences>

Turnitin (<http://www.turnitinuk.com/>) is an Internet-based 'originality checking tool' which allows documents to be compared with content on the Internet, in journals and in an archive of previously submitted works. It can help to detect unintentional or deliberate plagiarism.

It is a formative tool that makes it easy for students to review their citations and referencing as an aid to learning good academic practice. Turnitin produces an 'originality report' to help guide you. To learn more about Turnitin go to:

https://guides.turnitin.com/01_Manuals_and_Guides/Student/Student_User_Manual